

AWS

Amazon web services

CONTENTS

Chapter	Topics
Global Infrastructure	A Conceptual Introduction to Amazon Web Services (AWS) Regions Availability Zones End points Creating an AWS Account Usage Tracking Billing
IAM	IAM Essentials IAM Policies IAM Users IAM Groups Securing your account – Creating an IAM user and setting up the CLI IAM Roles
Cloud Watch	Alarms Events Monitoring Logging
AWS Simple Notification Service SNS	Introduction to SNS Understudying Topics Subscription and Publisher concepts
CloudTrail	Understanding the concepts CloudTrail
VPC	Basic understanding of VPC Subnets Route Table Internet Gateway NAT Gateway Security in VPC: NACL VPC Peering VPC Endpoints Limitation of VPC
Server Based Compute (EC2) Fundamentals	EC2 Architecture Instance, Types and Sizes EC2 instance purchasing types Instance Roles AMI's Ec2 Storage Architecture EBS, Types EFS Overview of Different types of storage classes Snapshots Life cycle manager Security Groups Elastic IPs Key pairs

	Load Balancing Autoscaling Route 53
--	---

Chapter	Topics
S3	Understanding S3 S3 Naming Convention Transferring Data to S3 Storage classes in S3 Lifecycle policy in S3 Cost optimization for S3 Versioning in S3 Encryption in S3 Static website and CORS

(1)
GLOBAL INFRASTRUCTURE

1. Aws Free tier account and monitoring

Billing dashboard

go to preferences and enable the receive free tier usage alert

2. Global infrastructure

We can access AWS i) AWS console ii) AWS cli iii) SDKs

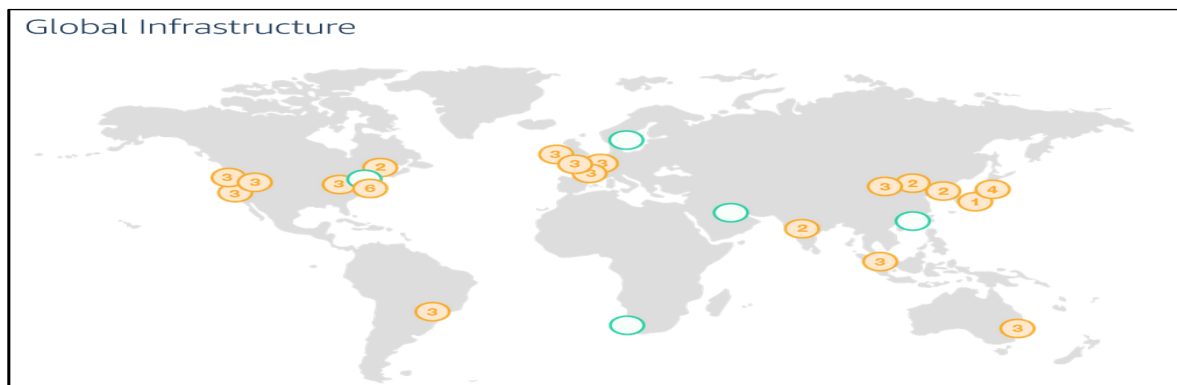
3. AWS Region:

- *AWS Region* is a separate geographic area where we cluster data centers.
- Each AWS Region is completely independent.

Availability zones: Each **AWS Region** consists of multiple, isolated, and physically separate AZ's within a geographic area.

Local zone: A *Local Zone* is an extension of an AWS Region that is geographically close to your users, for low latency communication.

Edge location: its data center, where end users access services **located** at AWS to reduce latency



4. Compute section:

EC2 Elastic compute cloud: Virtual machine

ECS Elastic Container service: Container as a service.

EBS Elastic beanstalk: PASS solution manages infrastructure for you.

AWS Lambda: Serverless computing platform.

5. Storage overview:

Databases: SQL Database - RDS, NOSQL Database - Dynamodb, Elastic cache, Redshift.

Storage: AWS S3 object storage service file used in application storing file and retrieving files.

(2)

IAM (Identity and Access Management)

- Provides access to accounts services where we can manage User, Roles, Groups & Policy password policy.
- It applies globally to all AWS regions.

Users: we create users and assign necessary permissions to them in the form of policies.

Groups: We can create groups for ex. Dev QA etc. and attach policies at the group level.

Policy:

A policy is a set of permission

Always explicit deny overrides explicit allow

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": "*",  
      "Resource": "*"   
    },  
  ],  
}
```

Policy types:

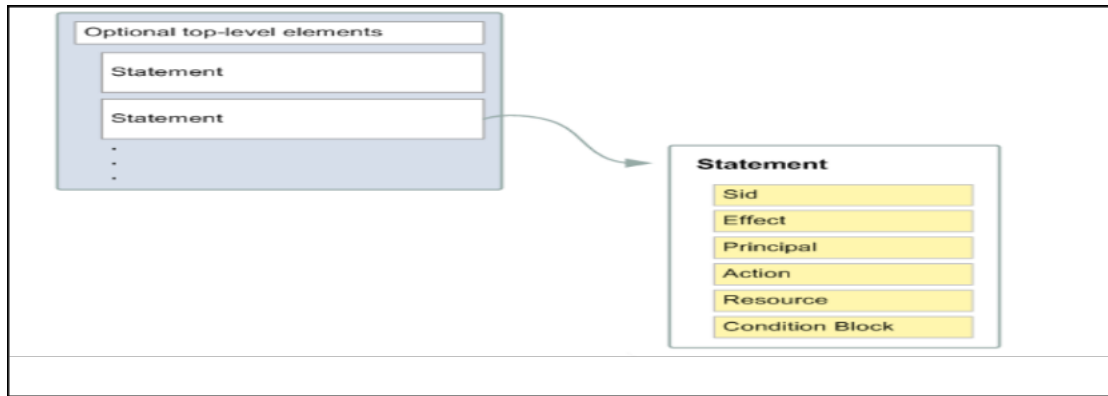
1. **Identity Based Policy:** Applicable on users, groups of users, and roles
 - AWS Managed policy:
 - Custom Managed Policy:
 - Inline Policy
2. **Resource Based policy:** Attach to a resource such as an Amazon S3 bucket
3. **Session based Policy:** create a temporary session for a role or federated user

Imp Notes:

More than one policy can be attached to a user or a group at the same time.

Policies can't be attached directly to resources like EC2 instance, S3 bucket etc.,

Basic Policy structure:



Effect : Can take only two value allow or deny

Principal: who is assuming the policy

Resource: on whom you are assuming the policy

Ex:

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "FirstStatement",
      "Effect": "Allow",
      "Action": ["iam:ChangePassword"],
      "Resource": "*"
    },
    {
      "Sid": "SecondStatement",
      "Effect": "Allow",
      "Action": "s3:ListAllMyBuckets",
      "Resource": "*"
    }
  ]
}
  
```

Q: Can we attach multiple policies to user group or a role.

Yes, we attach.

Roles:

- Roles comes between services like ec2 wants to access S3 or non-AWS user (hybrid account) should access AWS Resources.
- Policies can't be attached to aws resources hence roles come into picture
- EC2 can be attached one role at a time.
- Can we assign multiple roles to a EC2 instance? No, we can't. we can assign only single role to EC2 instance.

Q: If an ec2 instance is not able to access s3 bucket what could be the reason

A Role needs to be attached with proper policy defined.

Assume Role:

Returns a set of temporary security credentials that you can use to access AWS resources that you might not normally have access to. These temporary credentials consist of an access key ID, a secret access key, and a security token.

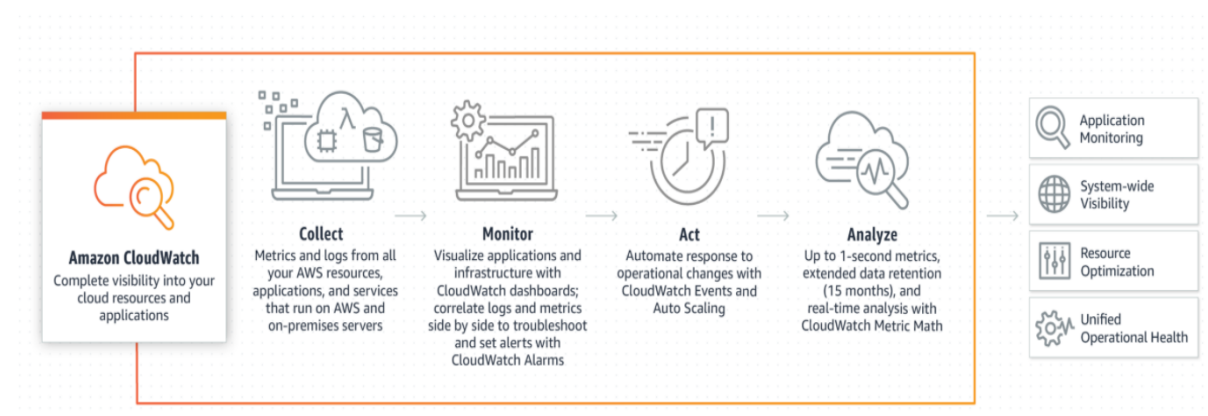
Example:

1. `aws s3api list-buckets`
2. Remove s3 access policy from the user
3. Try to list the bucket - you can't list the bucket because the user is not having the permission to list bucket (access denied)
4. Grant the user to assume a role
 - Create a role
 - Attach a policy s3 full access policy
 - get inside the role --- trust Relationship change it to user instead of ec2
 - AWS: "ARN OF IAM USER"
5. `aws sts assume-role --role-arn <enter the role arn> --role-session-name s3-access-example`
6. copy the accesskey secretkey and session token
- 7 `set AWS_ACCESS_KEY_ID=<enter the copied access key>`
`set AWS_SECRET_ACCESS_KEY=<enter the copied session key>`
`set AWS_SESSION_TOKEN=<sessiontoken>`
8. `aws s3api list-buckets`
9. Remove the env variable
`set AWS_ACCESS_KEY_ID=`
`set AWS_SECRET_ACCESS_KEY=`
`set AWS_SESSION_TOKEN=`

The AWS Security Token Service (STS) is a web service that enables you to request temporary, limited-privilege credentials for AWS Identity and Access Management (IAM) users or for users that you authenticate (federated users).

(3) CloudWatch

Amazon CloudWatch monitors your Amazon Web Services (AWS) resources and the applications you run on AWS in real time



Default metrics of EC2 instance: Network usage CPU Usage

Metrics:

Metrics are data about the performance of your systems

Basic monitoring: which polls for every 5 minutes

Detailed monitoring: which polls for every 1 minute.

Alarm:

If CPU utilization goes beyond the static threshold alarm goes to alarm state

Alarm state

Insufficient

OK state

Assignment: set a billing alarm to your account

Events: An Event indicates change in AWS environment

Event Resource: Which resource you want to monitor
Event target: to alert the event change through notifications

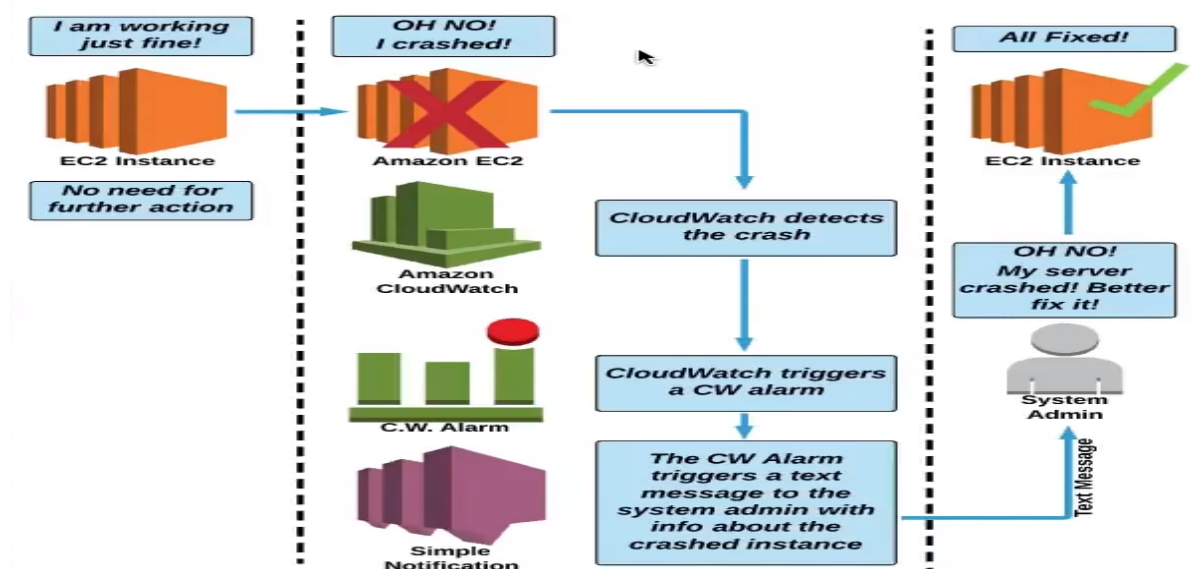
Logs:

CloudWatch Logs enables you to centralize the logs from all your systems, applications, and AWS services

Assignment follow the link and try to reproduce the same:

Ex: <https://www.youtube.com/watch?v=F4IE69V-iuw>

SNS Usage Example/Workflow:



(4)

Simple Notification Service

- Amazon Simple Notification Service is a notification service provided as part of Amazon Web Service.
- It provides a low-cost infrastructure for the mass delivery of messages, predominantly to mobile users

Topic:

An Amazon SNS topic is a logical access point that acts as a communication channel



(5) CloudTrail

- Auditing tool records all was account activity.
- Any action taken by users, roles and AWS services are recorded to cloud trial.
- Cloud trial events are kept for 90 days in event history
- You can create a trail of your own store the event history in s3 bucket.
- There are two types of event

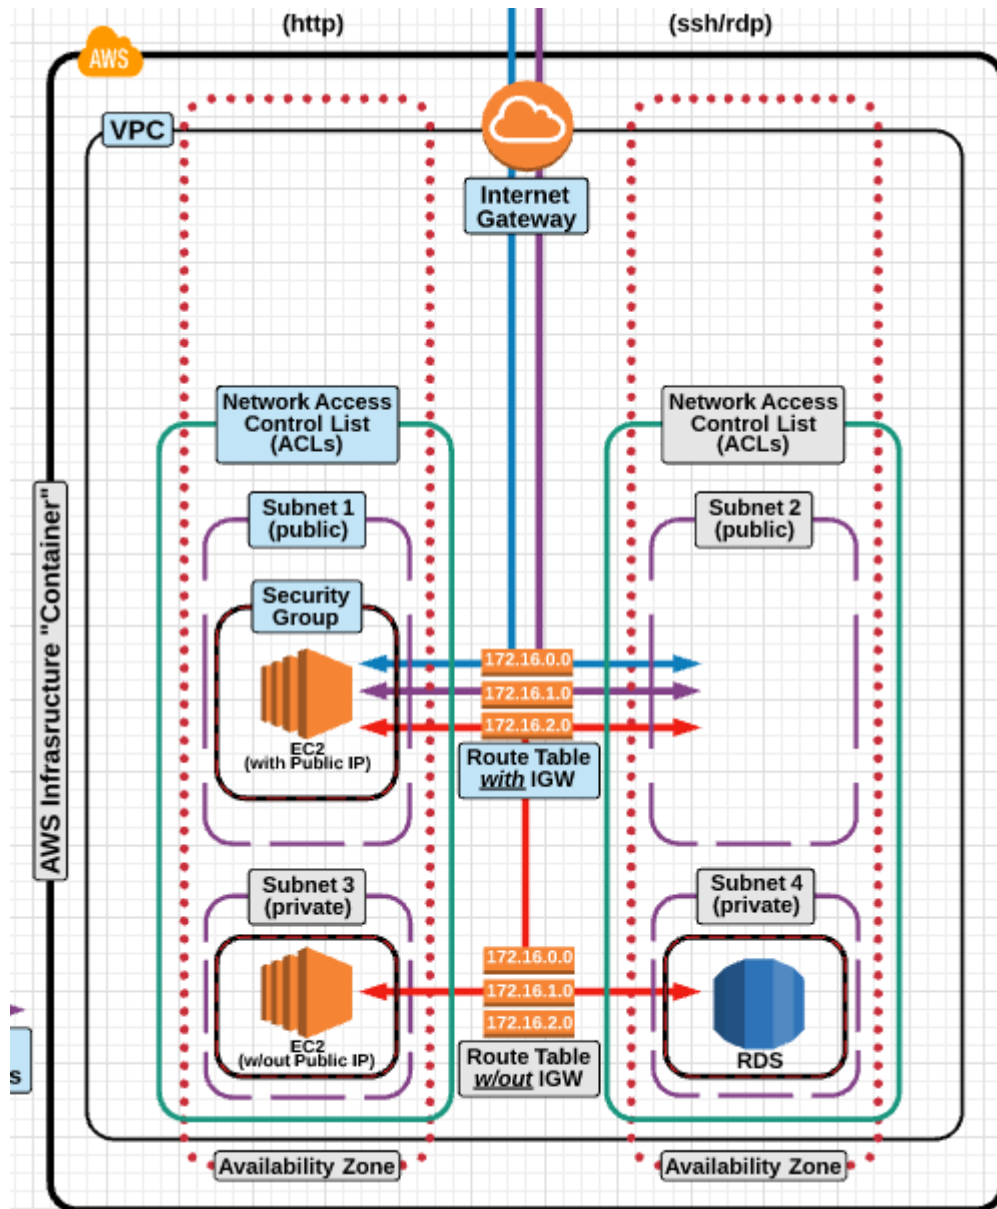
Management events: Management operations performed on AWS

Data events : currently supported S3 and Lambda: You can now record all API actions on S3 Objects and receive detailed information such as the AWS account of the caller, IAM user role of the caller, time of the API call, IP address of the API, and other details

(6)

Virtual Private Cloud

Amazon Virtual Private Cloud (Amazon VPC) lets you provision a logically isolated section of the AWS Cloud where you can launch AWS resources in a virtual network that you define like EC2 instance Databases.



CIDR:(Classes interdomain routing)

Classless Inter-Domain Routing is a method for allocating IP addresses and for IP routing.

Ex: The IPv4 block 192.168.100.0/22 represents the 1024 IPv4 addresses from 192.168.100.0 to 192.168.103.255.

I.e. $2^{(32-22)} = 2^{10} = 1024$ IPv4 addresses.

VPC design:

VPC CIDR = 10.180.0.0/16 means we have 65536 IPv4 address			
IPv4 Address range is 10.180.0.0 ---- 10.180.255.255			
Public subnet 1	Public Subnet 2	Private Subnet 1	Private Subnet 2
10.180.0.0/24	10.180.1.0/24	10.180.2.0/24	10.180.3.0/24
256 IPV4 address	256 IPV4 address	256 IPV4 address	256 IPV4 address
10.180.0.0-10.180.0.255	10.180.1.0-10.180.1.255	10.180.2.0– 10.180.2.155	10.180.3.0 – 10.180.3.255

- A **subnet** is a range of IP addresses in your VPC. You can launch AWS resources into a specified subnet.
- Some IP addresses are reserved they are
 - 10.180.0.0 Network address
 - 10.180.0.1 VPC Router
 - 10.180.0.2 DNS server (**DNS**. (Domain Name System) The Internet's system for converting alphabetic names into numeric IP addresses)
 - 10.180.0.3 Future use
 - 10.180.0.4 N/W Broadcast address
- VPC spans multiple Availability zones.
- Subnets must be associated with route table
 - A public subnet has a route to internet
 - A private subnet doesn't have route to internet. It creates higher level of security.
- You can use a network address translation (NAT) gateway to enable instances in a private subnet to connect to the internet or other AWS services, but prevent the internet from initiating a connection with those instances.

VPC Quota or VPC limitations

- 5 VPC per region
- 5 IGW per region
- Subnet per VPC 200
- IPv4 CIDR blocks per VPC 4
- Elastic IP addresses per Region 5
- Internet gateways per Region 5

- NAT gateways per Availability Zone 5
- Network ACLs per VPC 200
- Rules per network ACL 200


VPC Peering:

- A VPC peering connection is a networking connection between two VPCs that enables you to route traffic between them using private IPv4 addresses or IPv6 addresses.
- Instances in either VPC can communicate with each other as if they are within the same network.
- You can create a VPC peering connection between your own VPCs, or with a VPC in another AWS account.
- The VPCs can be in different regions (also known as an inter-region VPC peering connection).



Conditions:

- CIDR block shouldn't overlap
- Transitive peering relationships are not supported. i.e here VPC B cannot connect with VPC C.
- If the VPCs are in different regions, inter-region data transfer costs apply.
- You cannot have more than one VPC peering connection between the same two VPCs at the same time.



NACL:

1. A network access control list (ACL) is an optional layer of security for your VPC that acts as a firewall for controlling traffic in and out of one or more subnets. (Firewall at subnet level)
 1. Inbound means – incoming (Ingress)

2. Outbound means – outgoing (egress)
3. Always explicit deny take precedence over allow

Rule #	Type	Protocol	Port Range	Source	Allow / Deny
80	SSH (22)	TCP (6)	22	0.0.0.0/0	ALLOW
90	HTTP (80)	TCP (6)	80	0.0.0.0/0	ALLOW
*	ALL Traffic	ALL	ALL	0.0.0.0/0	DENY

Security Group:

A **security group** acts as a virtual firewall for your instance to control inbound and outbound traffic.

Security group rules:

Inbound

Outbound

Type ⓘ	Protocol ⓘ	Port Range ⓘ	Source ⓘ	Description ⓘ
SSH ▾	TCP	22	Anywhere ▾ 0.0.0.0/0, ::/0	Admin access.
HTTP ▾	TCP	80	Anywhere ▾ 0.0.0.0/0, ::/0	Web traffic.
HTTPS ▾	TCP	443	Custom ▾ 0.0.0.0/0, ::/0	Secure web traffic.

Add Rule

(6)

Elastic Cloud Compute

An **EC2 instance** is a virtual server in Amazon's Elastic Compute Cloud (**EC2**)

EC2 instance types:

EC2 Instance Type:

- Instance Types describe the "hardware" components that an EC2 instance will run on:
 - Compute power (processor/vCPU)
 - Memory (ram)
 - Storage Options/optimization (hard drive)
 - Network Performance (bandwidth)
- Instance Types are grouped into families and types that you can choose from that have different purposes:
 - General Purpose (T2, M5, and M4):
 - T2 - Burstable performance, good for many general purposes
 - M4/M5 - Small or mid-size databases, data processing, enterprise applications
 - Compute Optimized - (C4 and C5):
 - High performance web servers, science/engineering apps, ad serving
 - Memory Optimized - (X1e, X1, and R4):
 - High performance databases, in-memory databases, large data processing engines
 - Accelerated Computing - (P3, P2, G3, F1):
 - P2/P3 - Machine/Deep learning, high performance databases, server-size GPU compute workloads
 - G3 - 3D visualizations and rendering, application streaming, video encoding, server-side graphics workloads
 - F1 - Genomics research, financial analytics, big data, and security
 - Storage Optimized - (H1, I3, D2):
 - D2/H1 - MapReduce, HDFS, network file systems, or data processing applications
 - I3 - NoSQL databases (Cassandra/MongoDB/Redis), data warehouses, Elasticsearch

Note that the instance families and types are the 'current' generation as of April 2018.

EC2 purchasing options:

EC2 Purchasing Options:

On-Demand:

- On-demand purchasing lets you choose any **instance type** and provision/terminate it at any time
- Is the **most expensive** purchasing option
- Is the **most flexible** purchasing option
- You are only charged when the instance is **running** (and billed by the second)

Reserved Instances (RI):

- Reserved purchasing allows you to purchase an instance for a **set time period** of one or three years
- This allows for a **significant price discount** over using on-demand
- You can select to pay upfront, partial upfront, or none upfront
- Once you buy a reserved instance, you own it for the selected time period and are **responsible for the entire price** - regardless of how often you use it
- Purchases of AZ-specific RIs provide capacity reservation in that AZ. Regional RI purchases do not - so it is theoretically possible AWS will run out of capacity

Spot Instances:

- Spot pricing is a way for you to **"bid"** on an instance type, and only pay for and use that instance when the spot price is **equal to or below** your "bid" price
- This option allows Amazon to sell the use of **unused instances**, for short amounts of time, at a **substantial discount**
- **Spot prices fluctuate** based on supply and demand in the spot marketplace
- You are **charged per second (with conditions)**
- When you have an active bid, an instance is **provisioned for you when the spot price is equal to or less than your bid price**
- A provisioned instances **automatically terminate when the spot price is greater than your bid price**.
- Bid on unused EC2 instances for "non production applications"

Dedicated Hosts:

- A dedicated physical machine that you have full control over. This can help save money on license fees and meet certain regulatory compliances

EBS (Elastic block storage)

EC2 Elastic Block Store (EBS) Basics:

- EBS volumes are **persistent**, meaning that they can live beyond the life of the EC2 instance they are attached to
- EBS backed volumes are **network attached storage**, meaning they can be attached/detached to or from various EC2 instances
- However, they can only be attached to ONE EC2 instance at a time
- EBS volumes have the benefit of being backed up into a **snapshot** - which can later be restored into a new EBS volume
- By default, EBS volumes are replicated within the Availability Zone
- EBS volumes are usually mounted to the file system at /dev/sda1 or /dev/xvda

EBS Types:

EC2 Elastic Block Store Volumes:

General Purpose SSD:

- Use for dev/test environments and smaller DB instances
- Performance of 3 IOPS/GB of storage size (burstable with baseline performance)
- Volume size of 1GB to 16TB
- Considerations when using T2 instances with SSD root volumes (burstable vs. baseline performance)

Provisioned IOPS SSD:

- Used for mission critical applications that require sustained IOPS performance
- Large database workloads
- Volume size of 4GB to 16TB
- Performs at provisioned level and can provision up to 32,000 IOPS per volume

Throughput Optimized HDD and Cold HDD:

- Cheaper than SSD options, also less performant
- Cold HDD - Designed for less-frequent access
- Volume size of 500GB - 16 TB
- Cannot be a boot volume

EBS Magnetic (Previous Generation):

- Low storage cost
- Used for workloads where performance is not important or data is infrequently accessed
- Volume size of Min 1GB Max 1 TB

EFS:

Amazon Elastic file system is a regional service storing data within and across multiple Availability Zones (AZs) for high availability and durability

Difference between EBS v/s EFS v/s S3

AMAZON EBS	AMAZON EFS	AMAZON S3
Hardly scalable	Scalable	Scalable
Block Storage	Object storage	Object Storage
Faster than S3 and EFS	Faster than S3, slower than EBS	Slower than EBS and EFS
Accessible only via the given EC2 Machine	Accessible via several EC2 machines and AWS services	Can be publicly accessible
Is meant to be EC2 drive	Good for shareable applications and workloads	Good for storing backups
File System interface	Web and file system interface	Web interface

Snapshot EBS

- You can back up the data on your Amazon EBS volumes to Amazon S3 by taking point-in-time snapshots. Snapshots are *incremental* backups

Snapshots are stored in S3

- Launch two ec2 instance in different az's(instance1 & instance2)
- Create EBS volume and attach it to instance1
- The volumes are attached to instance1 you can verify it by logging into instance1 and executing "lsblk" command, but it's not mounted you can verify it through by running command "df -TH"

- Mount the volume to instance1
 - Format the disk with ext4: "mkfs -t ext4 /dev/xvdf"
 - Create a directory in root: 1. "cd /" 2. "mkdir /mnt/mydisk"
 - Mount the disk: "mount /dev/xvdf /mnt/mydisk"
 - you can verify that disk is mounted by running "df -TH" command.
- Create some files
- Take a snapshot
- Unmount the disk
 - umount /mnt/mydisk
 - Detach the volume from ec2 instance.
 - delete the volume
- Create a new volume from snapshot
- Attach the volume to newly created instance2.
- Mount the volume to instance2
 - Create a directory in root: 1. cd / 2. mkdir /mnt/mydisk
 - mount /dev/xvdf /mnt/mydisk

Assignment: Difference between instance level snapshot and volume snapshot

Data life cycle Manager:

You can use Amazon Data Lifecycle Manager to automate the creation, retention, and deletion of snapshots taken to back up your Amazon EBS volumes

Amazon machine image (AMI):

- An Amazon Machine Image (AMI) provides the information required to launch an instance.
- You can launch multiple instances from a single AMI when you need multiple instances with the same configuration.


Difference between Snapshot and AMI

An EBS snapshot is a backup of a single EBS volume. The EBS snapshot contains all the data stored on the EBS volume at the time the EBS snapshot was created.

An AMI image is a backup of an entire EC2 instance. Associated with an AMI image is EBS snapshots. Those EBS snapshots are the backups of the individual EBS volumes attached to the EC2 instance at the time the AMI image was created.

Elastic load Balancer (ELB):

Manage and control the flow of inbound request to group of targets by distributing the requests evenly across the targets. The targets may be EC2 instances lambda or containers.



Types of Load balancer:

Application load balancer:

- Used mainly for web application running http and https protocols.
- Operates at request level.

Network Load balancer:

- Ultra-high Performance at very low latency.
- Operates at connection level, routing traffic to targets with in VPC.
- Can handle millions of requests per second.

Classic load Balancer:

- Used for applications that were built in existing EC2 classic env.
- Operates both at connection & request level.

Example: Classic load balancer


- Spin up an EC2 instance1 in another availability zone (az1) with http port open in Security group
- Add the below script and launch the instance.


```
#!/bin/bash
yum update -y
yum install httpd -y
echo '<h1> Resonse from server-1 </h1>' > /var/www/html/index.html
service httpd start
chkconfig httpd on
```
- Spin up one more EC2 instance1 in another availability zone (az2) with http port open in Security group
- Add the below script and launch the instance.


```
#!/bin/bash
yum update -y
yum install httpd -y
echo '<h1> Resonse from server-2 </h1>' > /var/www/html/index.html
service httpd start
chkconfig httpd on
```

Autoscaling

AWS Auto Scaling monitors your applications and automatically adjusts capacity to maintain steady, predictable performance at the lowest possible cost.

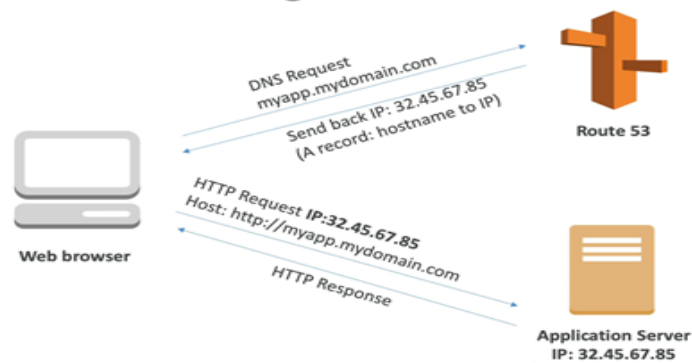


(7)
Route53:

Amazon Route 53 is a highly available and scalable Domain Name System (DNS) web service, where we can point IP address to domain name or point host name to another host

name.

Route 53 – Diagram for A Record



A Record: Maps IP address to domain name ex: 10.180.0.0 to myapp.mydomain.com

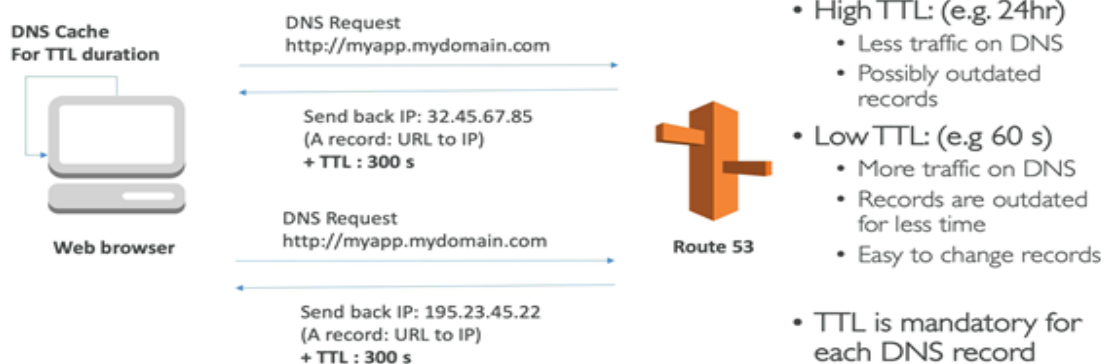
The screenshot displays the AWS Route 53 console. On the left is a navigation menu with options like Dashboard, Hosted zones, Health checks, Traffic flow, Domains, Resolver, VPCs, and Inbound/outbound endpoints. The main area shows a table of record sets for the domain 'stephanetheteacher.com'.

Name	Type	Value	Evaluate Target Health	Health
stephanetheteacher.com.	A	104.198.14.52	-	-
stephanetheteacher.com.	NS	ns-1468.awsdns-55.org, ns-633.awsdns-16.net, ns-252.awsdns-31.com, ns-1800.awsdns-33.co.uk.	-	-
stephanetheteacher.com.	SOA	ns-252.awsdns-31.com. awsdns-hostmaster.amazon	-	-
cname1test.stephanetheteacher.com.	CNAME	web-app-alb-307564896.eu-west-3.elb.amazonaws.	-	-
myfirstrecord.stephanetheteacher.com.	A	11.22.33.44	-	-
www.stephanetheteacher.com.	CNAME	kind-booht-aa82ac.netlify.com	-	-

On the right, the 'Create Record Set' form is visible, showing fields for Name, Type (A - IPv4 address), TTL (Seconds) set to 30, and Value (34.255.122.72).

TTL

DNS Records TTL (Time to Live)



Dashboard
Hosted zones
Health checks
Traffic flow
Traffic policies
Policy records
Domains
Registered domains
Pending requests
Resolver
VPCs
Inbound endpoints
Outbound endpoints
Rules

Back to Hosted Zones Create Record Set Import Zone File Delete Record Set Test Record Set

Record Set Name: Any Type: Aliases Only Weighted Only

Displaying 1 to 6 out of 6 Record Sets

Name	Type	Value	Evaluate Target Health	Health
stephanetheteacher.com.	A	104.198.14.52	-	-
stephanetheteacher.com.	NS	ns-1468.awsdns-55.org. ns-633.awsdns-15.net. ns-252.awsdns-31.com. ns-1800.awsdns-33.co.uk.	-	-
stephanetheteacher.com.	SOA	ns-252.awsdns-31.com. awsdns-hostmaster.amazon	-	-
cname-test.stephanetheteacher.com.	CNAME	web-app-alb-307564896.eu-west-3.elb.amazonaws.	-	-
myfirstrecord.stephanetheteacher.com.	A	11.22.33.44	-	-
www.stephanetheteacher.com.	CNAME	kind-booth-a82ac.netlify.com	-	-

Create Record Set

Name: ttl-demo .stephanetheteacher.co

Type: A - IPv4 address

Alias: Yes No

TTL (Seconds): 300 +1m 5m 1h 1d

Value: 34.255.122.72

IPv4 address. Enter multiple addresses on separate lines. Example: 192.0.2.235 198.51.100.234

Routing Policy: Simple

Route 53 responds to queries based only on the values in this record. [Learn More](#)

CNAME Record:

Maps hostname to another host name: us-east.2.elb.amazonaws.com to myapp.mydomain.com

Dashboard
Hosted zones
Health checks
Traffic flow
Traffic policies
Policy records
Domains
Registered domains
Pending requests
Resolver
VPCs
Inbound endpoints
Outbound endpoints
Rules

Back to Hosted Zones Create Record Set Import Zone File Delete Record Set Test Record Set

Record Set Name: Any Type: Aliases Only Weighted Only

Displaying 1 to 6 out of 6 Record Sets

Name	Type	Value	Evaluate Target Health	Health
stephanetheteacher.com.	NS	ns-252.awsdns-31.com. ns-1468.awsdns-55.org. ns-633.awsdns-15.net. ns-1800.awsdns-33.co.uk.	-	-
stephanetheteacher.com.	SOA	ns-252.awsdns-31.com. awsdns-hostmaster.amazon	-	-
cname-test.stephanetheteacher.com.	CNAME	web-app-alb-307564896.eu-west-3.elb.amazonaws.	-	-
myfirstrecord.stephanetheteacher.com.	A	11.22.33.44	-	-
ttl-demo.stephanetheteacher.com.	A	3.86.116.186	-	-
www.stephanetheteacher.com.	CNAME	kind-booth-a82ac.netlify.com	-	-

Create Record Set

Name: myapp .stephanetheteacher.com

Type: CNAME - Canonical name

Alias: Yes No

TTL (Seconds): 300 1m 5m 1h 1d

Value: DemoRoute531326827296.eu-west-1.elb.amazonaws.com

The domain name that you want to resolve to instead of the value in the Name field. Example: www.example.com

Routing Policy: Simple


Route 53 responds to queries based only on the values in this record. [Learn More](#)

Create

Alias Record: points a host name to AWS Resource ex: myapp.mydomain.com to us-east.2.elb.amazonaws.com


Latency Routing Policy:

Use when you have resources in multiple AWS Regions, and you want to route traffic to the region that provides the best latency.



Weighted Routing Policy:

Use to route traffic to multiple resources in proportions that you specify.



(8)

Simple storage service (S3)

Amazon S3 has a simple web services interface that you can use to store and retrieve any amount of data, at any time, from anywhere on the web.

Single operation upload:

- It's a traditional upload where you will upload the object in one part
- A single operation upload can upload the file up to 5GB in size.

Upload object in parts:

- Using multipart upload, you can upload the large objects up to 5TB.
- You can use multipart upload for the objects from 5MB to 5TB in size.

Rules for bucket naming:

- Bucket names must be between 3 and 63 characters long.
- Bucket names can consist only of lowercase letters, numbers, dots . and hyphens -.
- Bucket names must begin and end with a letter or number.
- Bucket names must not be formatted as an IP address (for example, 192.168.5.4).
- Bucket names can't begin with xn-- (for buckets created after February 2020).

Limitation of S3 bucket:

- Only 100 buckets can be created per account.
- Can hold unlimited objects

S3 Storage classes:

- Standard:
 - Designed for general- and all-purpose storage
 - Default storage option
 - 99.999999999% object durability
 - 99.99% object availability
 - Most expensive storage class.
- Reduced Redundancy storage
 - Designed for non-critical objects
 - 99.99% object durability
 - 99.99% object availability
 - Less expensive than standard
- Infrequent access
 - Designed for less frequently accessed objects.
 - 99.999999999% object durability
 - 99.99% object availability

Less expensive than reduced redundancy storage

- Glacier
 - Designed for long term archival storage
 - May take several hours to retrieve the objects from this storage
 - Cheapest s3 storage class

S3 Life cycle policy:


An object lifecycle policy is a set of rules that automate the migration of the object storage class to different storage class

By default, lifecycle policies are disabled for a bucket

S3

Example:


- I have a work file that I am going to access every day for the next 30 days
- After 30 days, I may only need to access that file once a week for the 60 next days
- After which (90 days total) I will probably never access the file again but want to keep it just in case



Encryption:

Two ways of protecting information with S3


1. Server side/At rest:



2. In-transit/Client-side encryption:



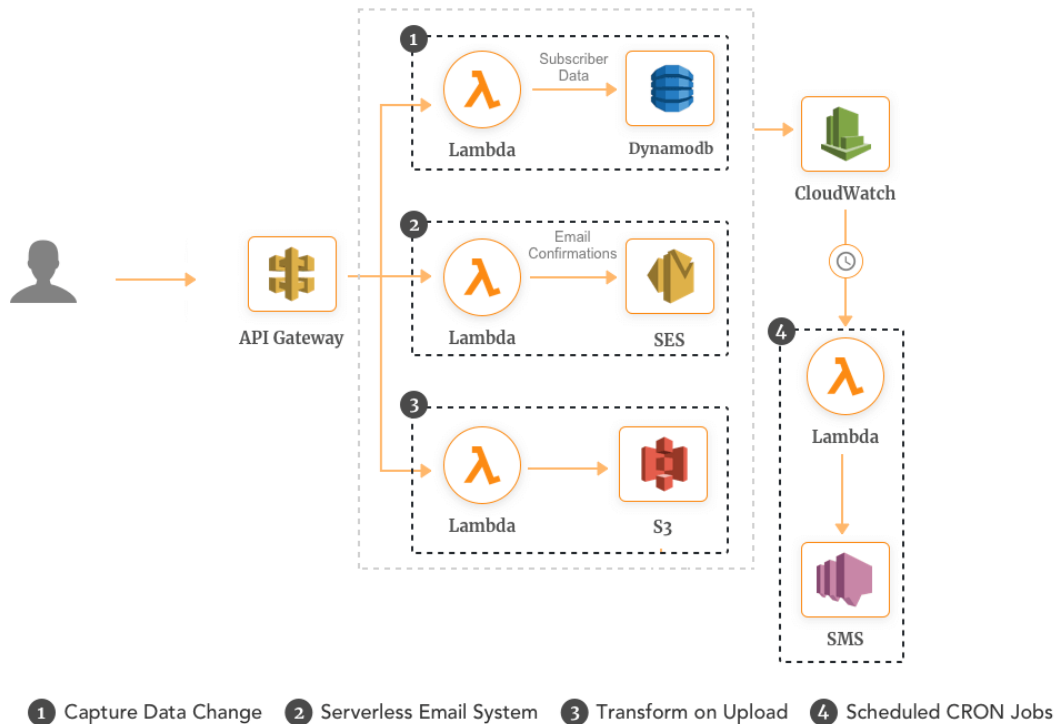
In client side/ in-transit we have two types



Lambda

AWS Lambda lets you run code without provisioning or managing servers. You pay only for the compute time you consume.

Just upload your code and Lambda takes care of everything required to run and scale your code with high availability



- Manage your virtual functions not really caring about the servers
- Run on demand
- Scaling is automated

Billing:

Pay per request first one million requests is free \$0.20 per one million request.
compute time 0.00001667 for every GB-seconds used.

AWS Lambda Languages:

NodeJS, Python, Python3, Gr00vy, java, csharp, Scala and GO

AWS Lambda Integration

Kinesis, API Gateway, DynamoDB, AWS S3, CloudWatch Events, CloudWatch logs, SNS and Incognito

Key Management Service (KMS)

AWS Key Management Service (AWS KMS) is a managed service that makes it easy for you to create and control customer master keys (CMKs), the encryption keys used to encrypt your data.

(11)

AWS secrets

AWS Secrets Manager helps you protect secrets needed to access your applications, services, and IT resources. The service enables you to easily rotate, manage, and retrieve database credentials, API keys, and other secrets throughout their lifecycle.

(12)

Relational Database Service (RDS)

Amazon Relational Database Service (Amazon RDS) makes it easy to set up, operate, and scale a relational database in the cloud.

Which relational database engines does Amazon RDS support?

Amazon RDS database engines:

- Amazon Aurora
- PostgreSQL
- MySQL
- MariaDB
- Oracle
- Microsoft SQL Server

Encryption in RDS:

Encryption at rest is supported for

- Amazon Aurora
- PostgreSQL
- MySQL
- MariaDB
- Oracle
- Microsoft SQL Server

Q: Can we enable encryption on existing DB

Encrypting existing DBs is not supported. To do this, you'll need to create a new encrypted instance and migrate data to it. The encryption key can be stored in KMS.

Q: Which is the non-relational database supported in AWS

Amazon DynamoDB is the NoSQL database supported by AWS

